

## ✓ HW 7 : GANs

In this assignment, you will learn to write an advanced PyTorch implementation concept commonly used in a complex deep learning pipeline by using GAN as a learning example.

You will also start working with more complex architectures (upsampling) and writing style (complex modules such as modulelist) for pytorch.

Every TODO is weighted equally. Optional TODO is half of a regular TODO.

```
1 !pip install torchsummary
```

```
● Collecting torchsummary
  Downloading torchsummary-1.5.1-py3-none-any.whl.metadata (296 bytes)
  Downloading torchsummary-1.5.1-py3-none-any.whl (2.8 kB)
  Installing collected packages: torchsummary
    Successfully installed torchsummary-1.5.1
```

```
1 import torchsummary
```

## ✓ GPU test

```
1 !nvidia-smi
```

```
Sat Apr  6 05:02:19 2024
+-----+
| NVIDIA-SMI 535.129.03      Driver Version: 535.129.03    CUDA Version: 12.2 |
|-----+-----+-----+-----+-----+-----+-----+-----+
| GPU  Name     Persistence-M | Bus-Id     Disp.A  | Volatile Uncorr. ECC | | | |
| Fan  Temp     Perf            Pwr:Usage/Cap | Memory-Usage | GPU-Util  Compute M. |
|          |              |             |              |          |          MIG M. |
|-----+-----+-----+-----+-----+-----+-----+-----+
|  0  Tesla P100-PCIE-16GB     Off | 00000000:00:04.0 Off |          0 | | | |
| N/A   38C   P0            27W / 250W |      0MiB / 16384MiB |     0%      Default |
|          |              |             |              |          |          N/A |
+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+
| Processes:                               GPU Memory |
| GPU  GI  CI      PID  Type      Process name        Usage |
| ID  ID
|-----+-----+-----+-----+-----+-----+-----+
| No running processes found
+-----+
```

## ✓ Part 1 : WGAN-GP reimplementaion

In this section, you are going to reimplement WGAN-GP (<https://arxiv.org/pdf/1704.00028.pdf>) based on the pseudocode provided in the paper to generate MNIST digit characters. Some parts are intentionally modified to discourage straight copypasting from public repositories.

**Algorithm 1** WGAN with gradient penalty. We use default values of  $\lambda = 10$ ,  $n_{\text{critic}} = 5$ ,  $\alpha = 0.0001$ ,  $\beta_1 = 0$ ,  $\beta_2 = 0.9$ .

---

**Require:** The gradient penalty coefficient  $\lambda$ , the number of critic iterations per generator iteration  $n_{\text{critic}}$ , the batch size  $m$ , Adam hyperparameters  $\alpha, \beta_1, \beta_2$ .  
**Require:** initial critic parameters  $w_0$ , initial generator parameters  $\theta_0$ .

- 1: **while**  $\theta$  has not converged **do**
- 2:   **for**  $t = 1, \dots, n_{\text{critic}}$  **do**
- 3:     **for**  $i = 1, \dots, m$  **do**
- 4:       Sample real data  $\mathbf{x} \sim \mathbb{P}_r$ , latent variable  $\mathbf{z} \sim p(\mathbf{z})$ , a random number  $\epsilon \sim U[0, 1]$ .
- 5:        $\hat{\mathbf{x}} \leftarrow G_\theta(\mathbf{z})$
- 6:        $\tilde{\mathbf{x}} \leftarrow \epsilon\mathbf{x} + (1 - \epsilon)\hat{\mathbf{x}}$
- 7:        $L^{(i)} \leftarrow D_w(\tilde{\mathbf{x}}) - D_w(\mathbf{x}) + \lambda(\|\nabla_{\tilde{\mathbf{x}}}D_w(\tilde{\mathbf{x}})\|_2 - 1)^2$
- 8:     **end for**
- 9:      $w \leftarrow \text{Adam}(\nabla_w \frac{1}{m} \sum_{i=1}^m L^{(i)}, w, \alpha, \beta_1, \beta_2)$
- 10:  **end for**
- 11:  Sample a batch of latent variables  $\{\mathbf{z}^{(i)}\}_{i=1}^m \sim p(\mathbf{z})$ .
- 12:   $\theta \leftarrow \text{Adam}(\nabla_\theta \frac{1}{m} \sum_{i=1}^m -D_w(G_\theta(\mathbf{z})), \theta, \alpha, \beta_1, \beta_2)$
- 13: **end while**

---

The pseudocode could be organized into two main parts: discriminator optimization in line 2-10, and generator optimization in line 11-12.

The discriminator part consists of four steps:

- Line 4: data, and noise sampling with a batch size of  $m$
- Line 5-7: discriminator loss calculation
- Line 9: discriminator update
- Repeat line 4-9 for  $n_{\text{critic}}$  steps

After the discriminator is updated, the generator is then updated by performing two steps:

- Line 11: noise sampling
- Line 12: generator loss calculation and update

This part is divided into four subsections: network initialization, hyperparameter initialization, data preparation, and training loop. The detail for each part will be explained in the subsections.

## ✓ Downloading MNIST dataset

The MNIST dataset contains 60,000 training digit character image (0-9) at 28x28 resolution that are normalized to [0, 1]. Given the training images, your task is to generate new training images using WGAN-GP by learning from the training distribution.

```

1 import torchvision.datasets as datasets
2 import numpy as np
3
4 mnist_trainset = datasets.MNIST(root='./data', train=True, download=True, transform=None)
5 trainX = np.array(mnist_trainset.data[..., None]).transpose(0, 3, 1, 2) / 255
6 print("Dataset size : ", trainX.shape)

Downloading http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz to ./data/MNIST/raw/train-images-idx3-ubyte.gz
100%|██████████| 9912422/9912422 [00:00<00:00, 78154681.11it/s]
Extracting ./data/MNIST/raw/train-images-idx3-ubyte.gz to ./data/MNIST/raw

Downloading http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz to ./data/MNIST/raw/train-labels-idx1-ubyte.gz
100%|██████████| 28881/28881 [00:00<00:00, 49788612.34it/s]
Extracting ./data/MNIST/raw/train-labels-idx1-ubyte.gz to ./data/MNIST/raw

Downloading http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz to ./data/MNIST/raw/t10k-images-idx3-ubyte.gz

100%|██████████| 1648877/1648877 [00:00<00:00, 19450044.99it/s]
Extracting ./data/MNIST/raw/t10k-images-idx3-ubyte.gz to ./data/MNIST/raw

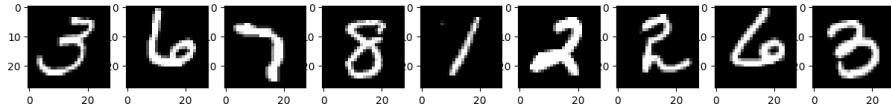
Downloading http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz to ./data/MNIST/raw/t10k-labels-idx1-ubyte.gz
100%|██████████| 4542/4542 [00:00<00:00, 18585881.72it/s]
Extracting ./data/MNIST/raw/t10k-labels-idx1-ubyte.gz to ./data/MNIST/raw

Dataset size : (60000, 1, 28, 28)

```

## Dataset Visualization

```
1 import matplotlib.pyplot as plt
2 import numpy as np
3 plt.figure(figsize = (15,75))
4 for i in range(9):
5     plt.subplot( int('19{}'.format(i+1)) )
6     plt.imshow( trainX[np.random.randint(len(trainX))].transpose((1, 2, 0))[..., 0] , cmap = 'gray' )
7 plt.show()
```



## Generator and Discriminator network

Before training, the deep learning networks have to be initialized first. Therefore, in this part, you are going to write a generator and discriminator network based on the description provided below.

The description of the discriminator network is shown in the Table below.

Discriminator (D(x))			
	Kernel size	Resample	Output shape
ConvBlock	$5 \times 5$	Down	$128 \times 14 \times 14$
ConvBlock	$5 \times 5$	Down	$256 \times 7 \times 7$
ConvBlock	$5 \times 5$	Down	$512 \times 4 \times 4$
Linear	-	-	1

The network also has some specific requirements:

- ConvBlock is a Convolution-ReLU layer
- All ReLUs in the encoder are leaky, with a slope of 0.1

The description of the generator network is shown in the Table below.

Generator (G(z))			
	Kernel size	Resample	Output shape
z	-	-	128
Linear	-	-	$512 \times 4 \times 4$
ConvBlock	$5 \times 5$	Up	$256 \times 8 \times 8$
ConvBlock	$5 \times 5$	Up	$128 \times 16 \times 16$
Conv, Sigmoid	$5 \times 5$	Up	$1 \times 32 \times 32$
-	-	Down	$1 \times 28 \times 28$

The network also has some specific requirements:

- ConvBlock is a ConvTranspose-BatchNorm-ReLU layer
- Downsampling method is bilinear interpolation (torch.nn.Upsample or torch.nn.functional.interpolate)

TODO 1: Implement a discriminator network.

TODO 2: Implement a generator network.

```

1 import torch
2 import torch.nn.functional as F
3 from torch import nn
4 from torchvision import transforms
5 from torchsummary import summary
6
7 class Discriminator(nn.Module):
8     ##TODO1 implement the discriminator (critic)
9     def __init__(self):
10         super().__init__()
11         self.conv_block1 = nn.Sequential(nn.Conv2d(1, 128, 5, 2, 2),
12                                         nn.LeakyReLU(0.1))
13         self.conv_block2 = nn.Sequential(nn.Conv2d(128, 256, 5, 2, 2),
14                                         nn.LeakyReLU(0.1))
15         self.conv_block3 = nn.Sequential(nn.Conv2d(256, 512, 5, 2, 2),
16                                         nn.LeakyReLU(0.1))
17         self.flatten = nn.Flatten()
18         self.output = nn.Linear(512*4*4, 1)
19
20     def forward(self, x):
21         x = self.conv_block1(x)
22         x = self.conv_block2(x)
23         x = self.conv_block3(x)
24         x = self.flatten(x)
25         x = self.output(x)
26
27     return x
28
29
30 class Generator(nn.Module):
31     ##TODO2 implement the generator (actor)
32     def __init__(self):
33         super().__init__()
34         self.input = nn.Linear(128, 512*4*4)
35         self.conv_block1 = nn.Sequential(nn.ConvTranspose2d(512, 256, 5),
36                                         nn.BatchNorm2d(256),
37                                         nn.ReLU())
38         self.conv_block2 = nn.Sequential(nn.ConvTranspose2d(256, 128, 5, 3, 5),
39                                         nn.BatchNorm2d(128),
40                                         nn.ReLU())
41         self.conv_block3 = nn.Sequential(nn.ConvTranspose2d(128, 1, 5, 3, 9),
42                                         nn.Sigmoid())
43
44     def forward(self, x):
45         x = self.input(x)
46         x = x.view(-1, 512, 4, 4)
47         x = self.conv_block1(x)
48         x = self.conv_block2(x)
49         x = self.conv_block3(x)
50         x = F.interpolate(x, (28, 28), mode='bilinear')
51
52     return x
53
54
55 discriminator = Discriminator().cuda()
56 generator = Generator().cuda()

```

## ✓ Network verification

TODO 3: What is the input and output shape of the generator and discriminator network? Verify that the implemented networks are the same as the answer you have provided.

```

1 data = torch.rand(10, 1, 28, 28).cuda()
2
3 summary(discriminator, (1, 28, 28))
4
5 discriminator(data).shape

```

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 128, 14, 14]	3,328
LeakyReLU-2	[-1, 128, 14, 14]	0

Conv2d-3	[-1, 256, 7, 7]	819,456
LeakyReLU-4	[-1, 256, 7, 7]	0
Conv2d-5	[-1, 512, 4, 4]	3,277,312
LeakyReLU-6	[-1, 512, 4, 4]	0
Flatten-7	[-1, 8192]	0
Linear-8	[-1, 1]	8,193

---

Total params: 4,108,289  
Trainable params: 4,108,289  
Non-trainable params: 0

---

Input size (MB): 0.00  
Forward/backward pass size (MB): 0.76  
Params size (MB): 15.67  
Estimated Total Size (MB): 16.44

---

torch.Size([10, 1])

```
1 data = torch.rand(10, 128).cuda()
```

```
2
```

```
3 summary(generator, (1, 128))
```

```
4
```

```
5 generator(data).shape
```

Layer (type)	Output Shape	Param #
<hr/>		
Linear-1	[-1, 1, 8192]	1,056,768
ConvTranspose2d-2	[-1, 256, 8, 8]	3,277,056
BatchNorm2d-3	[-1, 256, 8, 8]	512
ReLU-4	[-1, 256, 8, 8]	0
ConvTranspose2d-5	[-1, 128, 16, 16]	819,328
BatchNorm2d-6	[-1, 128, 16, 16]	256
ReLU-7	[-1, 128, 16, 16]	0
ConvTranspose2d-8	[-1, 1, 32, 32]	3,201
Sigmoid-9	[-1, 1, 32, 32]	0

---

Total params: 5,157,121  
Trainable params: 5,157,121  
Non-trainable params: 0

---

Input size (MB): 0.00  
Forward/backward pass size (MB): 1.20  
Params size (MB): 19.67  
Estimated Total Size (MB): 20.88

---

torch.Size([10, 1, 28, 28])

## Parameter Initialization

After the network is initialized, we then set up training hyperparameters for the training. In this part, hyperparameters have already been partially provided in the cell below, though some of them are intentionally left missing (None). Your task is to fill the missing parameters based on the pseudocode above.

TODO4: Initialize the missing model hyperparameters and optimizers based on the pseudocode above.

**Note:** To hasten the training process of our toy experiment, the training step and batch size is reduced to 3000 and 32, respectively.

```
1 NUM_ITERATION = 3000
2 BATCH_SIZE = 32
3 fixed_z = torch.randn((8, 128)).cuda()
4 def schedule(i):
5     lr = 1e-4
6     if(i > 2500): lr *= 0.1
7     return lr
8 losses = {'D' : [None], 'G' : [None]}
9
10 ## TODO4 initialize missing hyperparameter and optimizer
11 G_optimizer = torch.optim.Adam(generator.parameters(),
12                                 lr=1e-4,
13                                 betas=(0, 0.9))
14 D_optimizer = torch.optim.Adam(discriminator.parameters(),
15                                 lr=1e-4,
16                                 betas=(0, 0.9))
17 GP_lambda = 10
18 n_critic = 5
```

## ✓ Data preparation

TODO 5: Create a dataloader that could generate the data in line 4. The dataloader should return  $\mathbf{x}, \mathbf{z}, \epsilon$  with a batch size of BATCH\_SIZE

```
1 trainX.shape
2 (60000, 1, 28, 28)
3
4 # TODO5 implement dataloader
5 from torch.utils.data import Dataset, DataLoader
6
7 class WGAN_dataset(Dataset):
8     def __init__(self, trainX):
9         self.trainX = torch.tensor(trainX).cuda()
10        self.eps = torch.rand((len(self.trainX),1), requires_grad=True).cuda()
11        self.z = torch.rand((len(self.trainX), 128), requires_grad=True).cuda()
12
13    def __len__(self):
14        return len(self.trainX)
15
16    def __getitem__(self, idx):
17        return (self.trainX[idx], self.z[idx], self.eps[idx])
18
19 wgan_dataset = WGAN_dataset(trainX)
20 wgan_dataloader = DataLoader(wgan_dataset, batch_size=BATCH_SIZE, shuffle=True)
```

## ✓ Training loop

This section is the place where the training section starts. **It is highly recommended that you understand the pseudocode before performing the tasks below.** To train the WGAN-GP you have to perform the following tasks:

TODO6: Update the learning rate base on the provided scheduler.

TODO7: Sample the data from the dataloader (Line 4).

TODO8 : Calcualte the discriminator loss (Line 5-7).

- In the line 7 you have to implement the gradient penalty term  $\lambda(||(\nabla_{\hat{\mathbf{x}}} D_w(\hat{\mathbf{x}}))||_2 - 1)^2$ , which is a custom gradient. You may read <https://pytorch.org/docs/stable/generated/torch.autograd.grad.html> to find how custom gradient is implemented.
- HINT: Gradient norm calculation is still part of the computation graph.

TODO9: Update the discriminator loss (Line 9).

TODO10: Calculate and update the generator loss (Line 11-12).

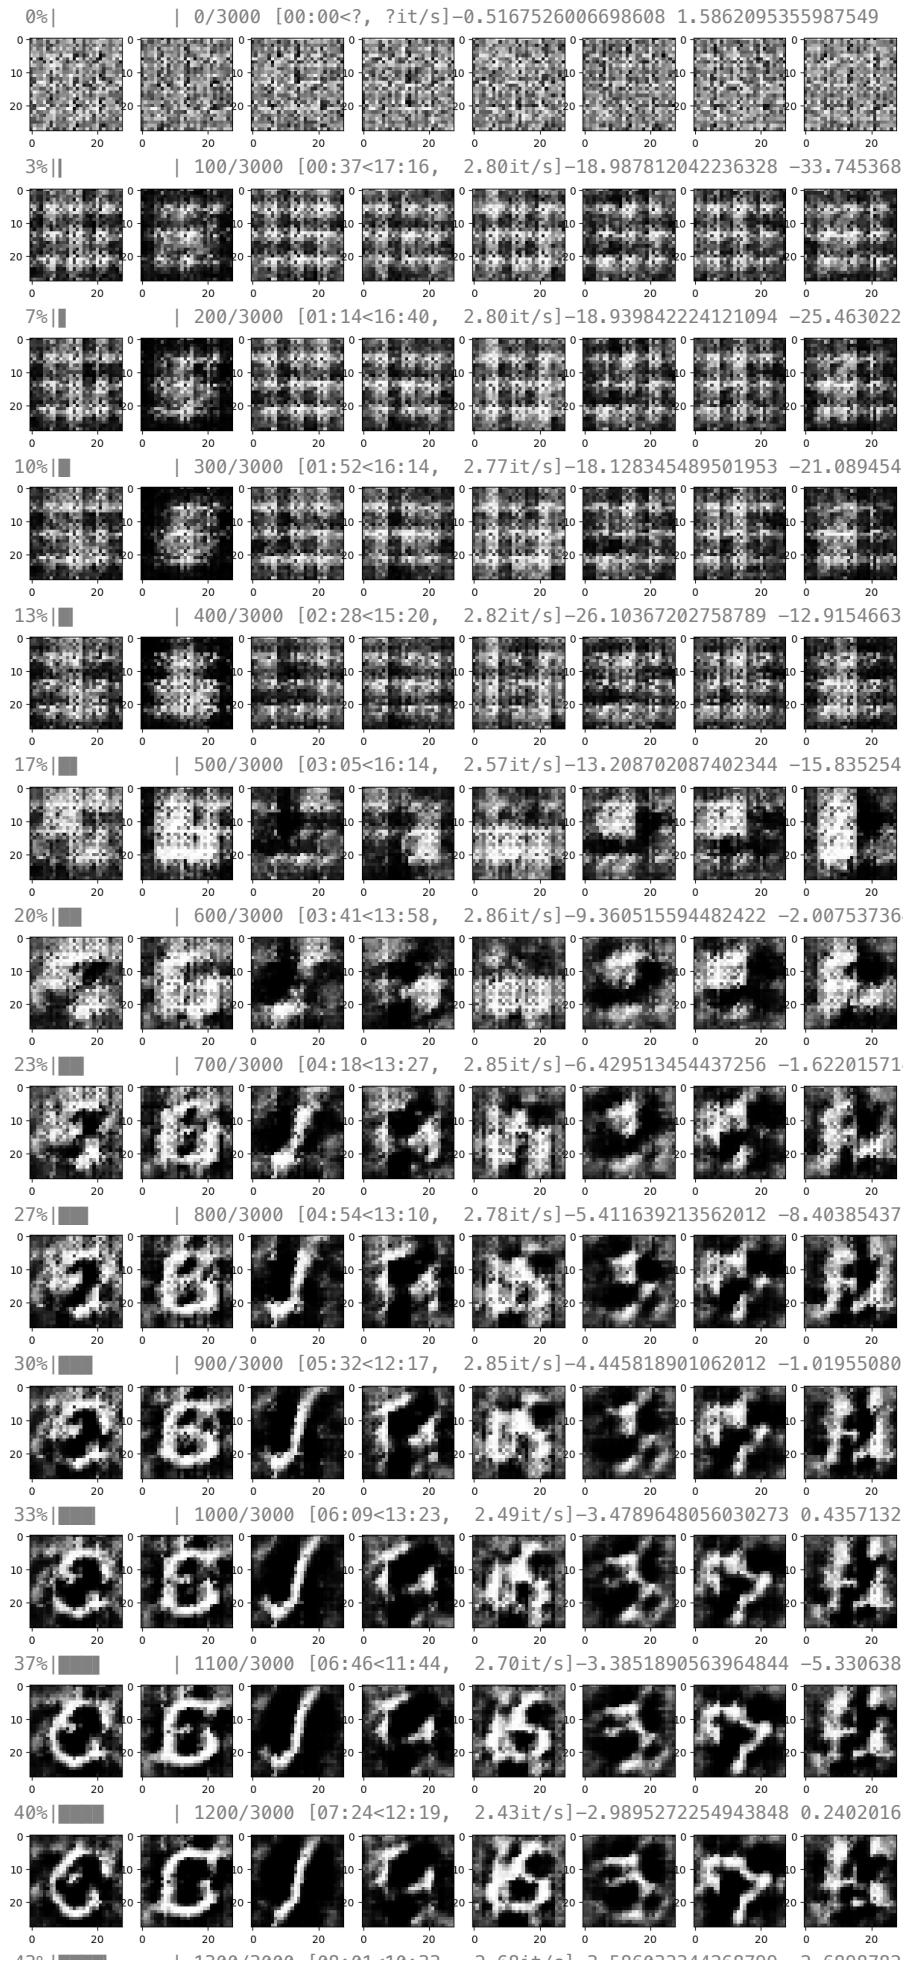
If your implementation is correct, the generated images should resemble an actual digit character after 500 iterations.

```
1 from torch import autograd
```

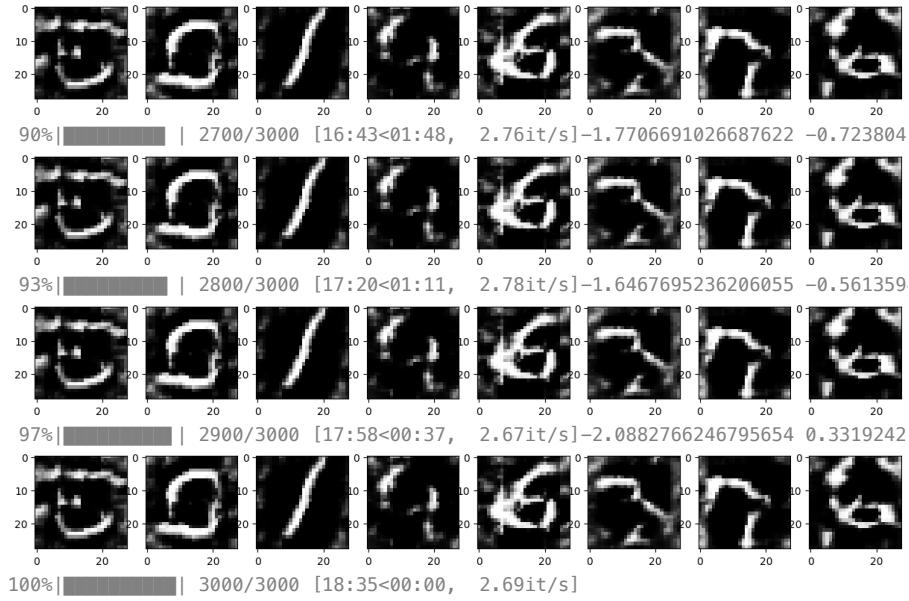
```

1 from tqdm import tqdm
2 for i in tqdm(range(NUM_ITERATION)):
3     ## TODO6 update learning rate
4     G_optimizer.param_groups[0]['lr'] = schedule(i)
5     D_optimizer.param_groups[0]['lr'] = schedule(i)
6
7     for t in range(n_critic):
8         generator.eval()
9         discriminator.train()
10
11    ## TODO7 line 4: sample data from dataloader
12    xs, zs, epss = next(iter(wgan_dataloader))
13    xs = xs.to(torch.float32)
14    epss = epss.unsqueeze(-1).unsqueeze(-1)
15
16    ## TODO8 line5-7 : calculate discriminator loss
17    x_tilde = generator(zs)
18    x_hat = epss*xs + (1-epss)*x_tilde
19    d_x_tilde = discriminator(x_tilde)
20    d_x = discriminator(xs)
21    d_x_hat = discriminator(x_hat)
22
23    grad_outputs = torch.ones_like(d_x_hat).cuda()
24
25    gradients = autograd.grad(
26        outputs=d_x_hat,
27        inputs=x_hat,
28        grad_outputs=grad_outputs,
29        create_graph=True,
30        retain_graph=True
31    )[0]
32
33    gradients = gradients.view(BATCH_SIZE, -1)
34    gradients = gradients.norm(2, 1)
35    gp = ((gradients - 1) ** 2)
36
37    loss = (d_x_tilde - d_x + gp).mean()
38    losses['D'].append(loss.item())
39
40    ## TODO9 : line 9 update discriminator loss
41    D_optimizer.zero_grad()
42    loss.backward()
43    D_optimizer.step()
44
45    ## TODO10 : line 11-12 calculate and update the generator loss
46    generator.train()
47    discriminator.eval()
48
49    _, zs, _ = next(iter(wgan_dataloader))
50    loss = -1*discriminator(generator(zs)).mean()
51    losses['G'].append(loss.item())
52
53    G_optimizer.zero_grad()
54    loss.backward()
55    G_optimizer.step()
56
57    # Output visualization : If your reimplementation is correct, the generated images should start resembling a digit character
58    if i % 100 == 0:
59        plt.figure(figsize = (15,75))
60        print(losses['D'][-1], losses['G'][-1])
61        with torch.no_grad():
62            res = generator(fixed_z).cpu().detach().numpy()
63            for k in range(8):
64                plt.subplot( int('18{}'.format(k+1)) )
65                plt.imshow( res[k].transpose(1, 2, 0)[..., 0], cmap = 'gray' )
66        plt.show()

```





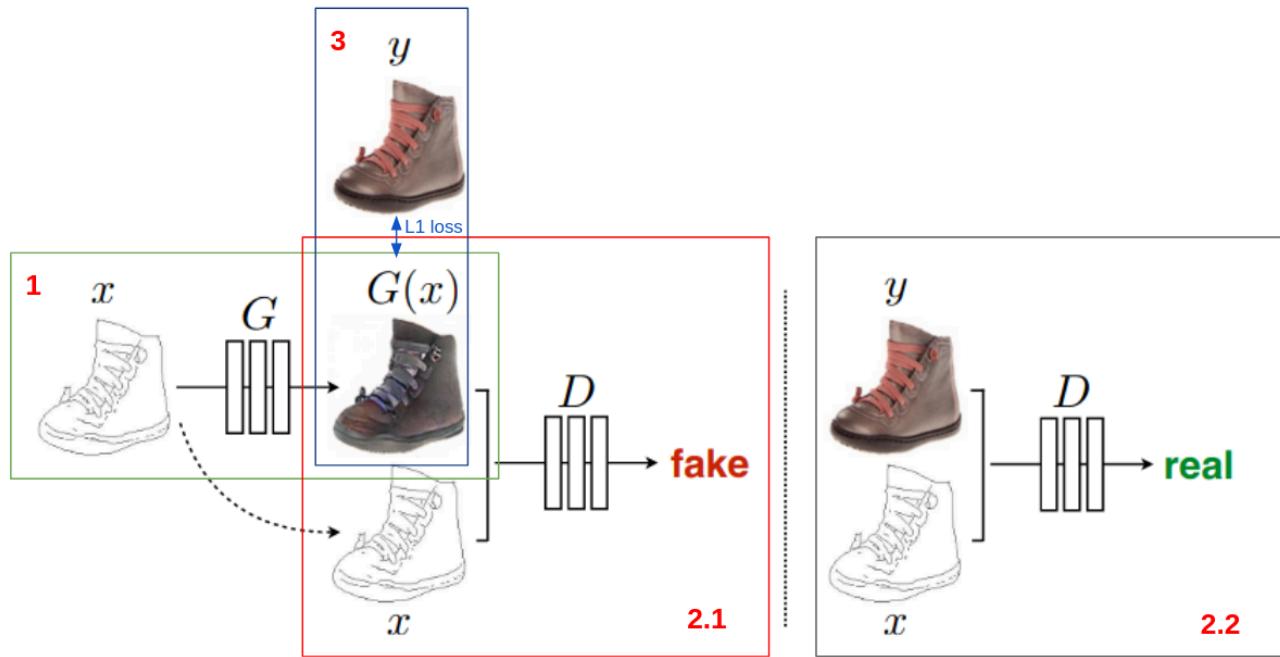


ref : <https://github.com/aadhithya/gan-zoo-pytorch/tree/master/models>

- ✓ Part 2 : pix2pix reimplementaion (cGAN on paired image translation)

In this exercise, we are reimplementing a paired image translation model, an application of a generative adversarial network (GAN). The model we are going to implement is pix2pix (<https://arxiv.org/pdf/1611.07004.pdf>), one of the earliest paired image translation models based on GAN.

The pipeline of pix2pix is shown in the Figure below.



From the figure above, the pipeline consists of three main parts:

- 1. Generation phase : the generator  $G$  create the generated image  $G(x)$  from the given input  $x$ .
- 2. Discrimination phase :

  - In step 2.1, the discriminator  $D$  receives an input image  $x$  and the generated image  $G(x)$ , then the discriminator has to learn to predict that the generated image  $G(x)$  is fake.
  - In step 2.2, the discriminator  $D$  receives an input image  $x$  and the ground truth image  $y$ , then the discriminator has to learn to predict that the image  $y$  is real.

- 3. Refinement phase: Refine the quality of the generated image  $G(x)$  by encouraging the generated image to be close to an actual image  $y$  by using L1 as an objective.

The objective of pix2pix is to train an optimal generator  $G^*$  base on the objective function :  $G^* = \text{argmin}_G \max_D L_{cGAN}(G, D) + \lambda L_1(G)$

- The term  $\text{argmin}_G \max_D L_{cGAN}(G, D)$  is the objective function of the first and second step, which is a standard cGAN loss :  $L_{cGAN}(G, D) = E_{x,y}[\log D(x, y)] + E_{x,z}[\log(1 - D(x, G(x, z)))]$ . The noise  $z$  is embedded in the generator in the form of dropout.
- The term  $L_1(G)$  is the objective function of the third step where  $L_1(G) = E_{x,y,z}[||y - G(x, z)||_1]$

The subsections will explain the dataset and training setup of this exercise.

## ✓ Get dataset

```

1 !wget http://efrosgans.eecs.berkeley.edu/pix2pix/datasets/facades.tar.gz
2 !tar -xzf facades.tar.gz

--2024-04-06 12:27:25-- http://efrosgans.eecs.berkeley.edu/pix2pix/datasets/facades.tar.gz
Resolving efrosgans.eecs.berkeley.edu (efrosgans.eecs.berkeley.edu)... 128.32.244.190
Connecting to efrosgans.eecs.berkeley.edu (efrosgans.eecs.berkeley.edu)|128.32.244.190|:80... connected.
HTTP request sent, awaiting response... 200 OK
Length: 30168306 (29M) [application/x-gzip]
Saving to: 'facades.tar.gz'

facades.tar.gz      100%[=====] 28.77M  472KB/s   in 72s

2024-04-06 12:28:37 (411 KB/s) - 'facades.tar.gz' saved [30168306/30168306]

```

## ✓ Import library

```

1 import cv2
2 import glob
3 import numpy as np
4 import torch
5 import torch.nn.functional as F
6 from torch import nn
7 from torchvision import transforms
8 import matplotlib.pyplot as plt
9 import matplotlib.gridspec as gridspec

```

## ✓ Setting up facade dataset

The dataset chosen for this exercise is the CMP Facade Database which is a pair of facade images and its segmented component stored in RGB value. The objective of this exercise is to generate a facade given its simplified segmented component. Both input and output is a 256 x 256 RGB image normalized to [-1, 1].

```

1 train = (np.array([cv2.imread(i) for i in glob.glob('facades/train/*')], dtype = np.float32) / 255).transpose((0, 3, 1, 2))
2 train = (train - 0.5) * 2 #shift from [0,1] to [-1, 1]
3 trainX = train[:, :, :, 256:]
4 trainY = train[:, :, :, :, :256]
5
6 val = (np.array([cv2.imread(i) for i in glob.glob('facades/val/*')], dtype = np.float32) / 255).transpose((0, 3, 1, 2))
7 val = (val - 0.5) * 2 #shift from [0,1] to [-1, 1]
8 valX = val[:, :, :, 256:]
9 valY = val[:, :, :, :, :256]
10
11 print("Input size : {}, Output size = {}".format(trainX.shape, trainY.shape))

Input size : (400, 3, 256, 256), Output size = (400, 3, 256, 256)

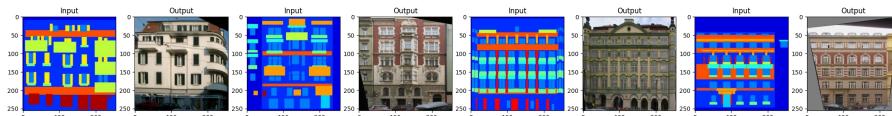
```

## ✓ Dataset Visualization

```

1 import matplotlib.pyplot as plt
2 import numpy as np
3 plt.figure(figsize = (30,90))
4 for i in range(4):
5     idx = np.random.randint(len(trainX))
6
7     plt.subplot( int('19{}'.format(2*i+1)) )
8     plt.title('Input')
9     plt.imshow( (0.5 * trainX[idx].transpose((1, 2, 0)) + 0.5)[..., ::-1] , cmap = 'gray' )
10    plt.subplot( int('19{}'.format(2*i+2)) )
11    plt.title('Output')
12    plt.imshow( (0.5 * trainY[idx].transpose((1, 2, 0)) + 0.5)[..., ::-1], cmap = 'gray' )
13 plt.show()

```

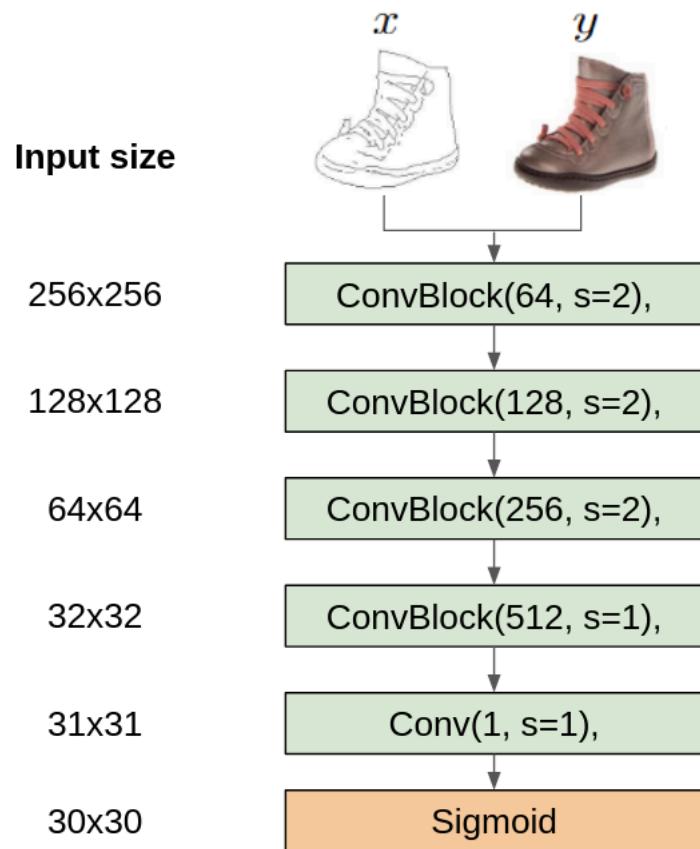


## Note

If you have trouble understanding the instruction provided in this homework or have any ambiguity about the instruction, you could also read the appendix section (section 6.1-6.2) in the paper for a detailed explanation.

## ✓ Discriminator network

In this section, we are going to implement a discriminator network of pix2pix. The description of the discriminator network is provided in the Figure below.



The network also has the following specific requirements:

- All convolutions are  $4 \times 4$  spatial filters
- ConvBlock is a Convolution-InstanceNorm-ReLU layer
- InstanceNorm is not applied to the first C64 layer
- All ReLUs are leaky, with a slope of 0.2

TODO 11: Implement the discriminator network based on the description above.

TODO 12: What should be the size of the input and output of the discriminator for this task? Verify that the input and output of the implemented network are the same as the answer you have provided.

```

1 class Discriminator(nn.Module):
2     #TOD011 implement the discriminator network
3     def __init__(self):
4         super().__init__()
5         self.ConvBlock1 = nn.Sequential(nn.Conv2d(6, 64, 4, stride=2, padding=1),
6                                         nn.LeakyReLU(0.2))
7         self.ConvBlock2 = nn.Sequential(nn.Conv2d(64, 128, 4, stride=2, padding=1),
8                                         nn.InstanceNorm2d(128),
9                                         nn.LeakyReLU(0.2))
10        self.ConvBlock3 = nn.Sequential(nn.Conv2d(128, 256, 4, stride=2, padding=1),
11                                         nn.InstanceNorm2d(256),
12                                         nn.LeakyReLU(0.2))
13        self.ConvBlock4 = nn.Sequential(nn.Conv2d(256, 512, 4, stride=1, padding=1),
14                                         nn.InstanceNorm2d(512),
15                                         nn.LeakyReLU(0.2))
16        self.ConvBlock5 = nn.Conv2d(512, 1, 4, stride=1, padding=1)
17        self.output = nn.Sequential(nn.Flatten(),
18                                     nn.LazyLinear(1),
19                                     nn.Sigmoid())
20
21    def forward(self, x):
22        x = self.ConvBlock1(x)
23        x = self.ConvBlock2(x)
24        x = self.ConvBlock3(x)
25        x = self.ConvBlock4(x)
26        x = self.ConvBlock5(x)
27        x = self.output(x)
28        return x
29
30 discriminator = Discriminator().cuda()

```

```

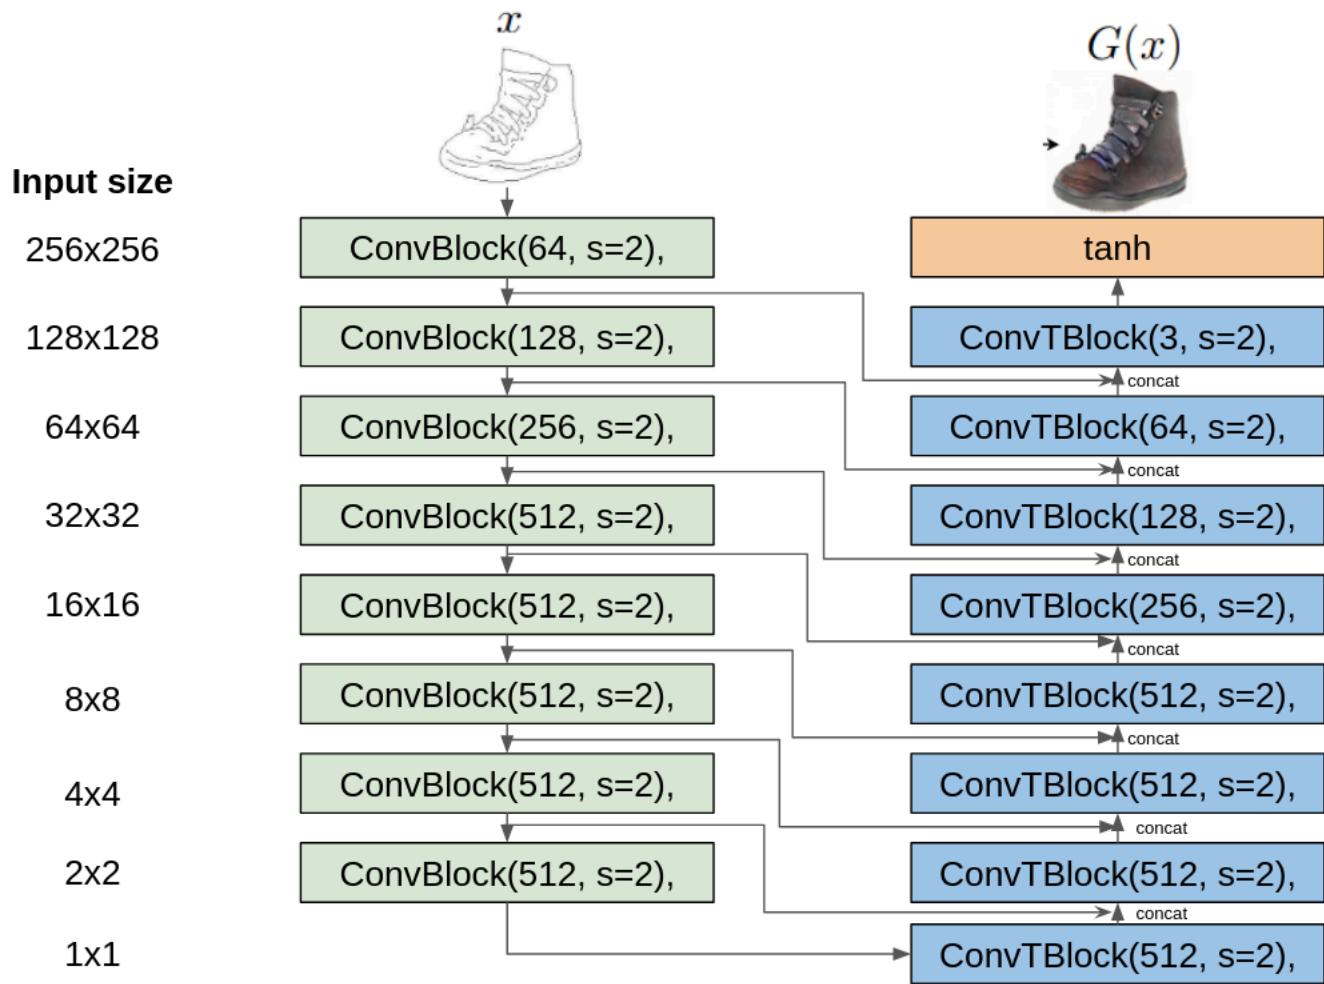
1 #TOD012 verify the discriminator
2 from torchsummary import summary
3 # input must concat in channel dimention (6, 256, 256)
4 # output must be prob of real or fake (1 to 0 respectively)
5 summary(discriminator, (6, 256, 256))

```

Layer (type)	Output Shape	Param #
<hr/>		
Conv2d-1	[−1, 64, 128, 128]	6,208
LeakyReLU-2	[−1, 64, 128, 128]	0
Conv2d-3	[−1, 128, 64, 64]	131,200
InstanceNorm2d-4	[−1, 128, 64, 64]	0
LeakyReLU-5	[−1, 128, 64, 64]	0
Conv2d-6	[−1, 256, 32, 32]	524,544
InstanceNorm2d-7	[−1, 256, 32, 32]	0
LeakyReLU-8	[−1, 256, 32, 32]	0
Conv2d-9	[−1, 512, 31, 31]	2,097,664
InstanceNorm2d-10	[−1, 512, 31, 31]	0
LeakyReLU-11	[−1, 512, 31, 31]	0
Conv2d-12	[−1, 1, 30, 30]	8,193
Flatten-13	[−1, 900]	0
Linear-14	[−1, 1]	901
Sigmoid-15	[−1, 1]	0
<hr/>		
Total params: 2,768,710		
Trainable params: 2,768,710		
Non-trainable params: 0		
<hr/>		
Input size (MB): 1.50		
Forward/backward pass size (MB): 45.28		
Params size (MB): 10.56		
Estimated Total Size (MB): 57.34		
<hr/>		

## ▼ Generator network

In this section, we are going to implement a generator network of pix2pix. The generator is based on the U-NET based architecture (<https://arxiv.org/abs/1505.04597>). The Description of the generator network is provided in the Figure below.



The network also has the following specific requirements:

- All convolutions are  $4 \times 4$  spatial filters
- ConvBlock is a Convolution-InstanceNorm-ReLU layer
- ConvTBlock is a ConvolutionTranspose-InstanceNorm-DropOut-ReLU layer with a dropout rate of 50%
- InstanceNorm is not applied to the first C64 layer in the encoder
- All ReLUs in the encoder are leaky, with a slope of 0.2, while ReLUs in the decoder are not leaky

TODO 13: Implement the generator network based on the description above.

TODO 14: What should be the size of the input and output of the generator for this task? Verify that the input and output of the implemented network are the same as the answer you have provided.

```

1 class Generator(nn.Module):
2     #TODO13 implement the generator network
3     def __init__(self):
4         super().__init__()
5         self.enc_layer_size = [(3, 64), (64, 128), (128, 256), (256, 512), (512, 512), (512, 512), (512, 512), (512, 512)]
6         self.enc_layers = nn.ModuleList([self.ConvBlock(i) for i in range(len(self.enc_layer_size))])
7
8         self.dec_layer_size = [(512, 512), (512*2, 512), (512*2, 512), (512*2, 512), (512*2, 256), (256*2, 128), (128*2, 64),
9         self.dec_layers = nn.ModuleList([self.ConvTBlock(i) for i in range(len(self.dec_layer_size))])
10
11         self.output = nn.Tanh()
12
13     def ConvBlock(self, i):
14         module = nn.Sequential()
15
16         module.append(nn.Conv2d(self.enc_layer_size[i][0], self.enc_layer_size[i][1], 4, stride=2, padding=1))
17         if i not in [0, 7]:
18             module.append(nn.InstanceNorm2d(self.enc_layer_size[i][1]))
19         module.append(nn.LeakyReLU(0.2))
20
21         return module
22
23     def ConvTBlock(self, i):
24         module = nn.Sequential()
25
26         module.append(nn.ConvTranspose2d(self.dec_layer_size[i][0], self.dec_layer_size[i][1], 4, stride=2, padding=1))
27         if i != 7:
28             module.append(nn.InstanceNorm2d(self.dec_layer_size[i][1]))
29             module.append(nn.Dropout(0.5))
30             module.append(nn.ReLU())
31
32         return module
33
34
35     def forward(self, x):
36         # encoder
37         encoder_track = [x]
38         for i in range(len(self.enc_layers)):
39             encoder_track.append(self.enc_layers[i](encoder_track[-1]))
40
41         # decoder
42         x = encoder_track[-1]
43         x = self.dec_layers[0](x)
44         for i in range(1, len(self.dec_layers)):
45             x = torch.cat([x, encoder_track[-1*(i+1)]], 1)
46             x = self.dec_layers[i](x)
47
48         # output
49         x = self.output(x)
50
51         return x
52
53 generator = Generator().cuda()

```

```

1 #HINT : you could also put multiple layers in a single list using nn.ModuleList
2 class Example(nn.Module):
3     def __init__(self):
4         super().__init__()
5         self.convs = nn.ModuleList([nn.Conv2d(25, 25, 3) for i in range(5)])
6     def forward(self, x):
7         for i in range(len(self.convs)):
8             x = self.convs[i](x)
9         return x
10 ex = Example().cuda()
11 print(ex(torch.zeros((8, 25, 32, 32)).cuda() ).shape)

torch.Size([8, 25, 22, 22])

```

```

1 #TODO14 verify the generator
2 summary(generator, (3, 256, 256))
3 generator(torch.zeros((32, 3, 256, 256)).cuda()).shape

```

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 64, 128, 128]	3,136

LeakyReLU-2	[-1, 64, 128, 128]	0
Conv2d-3	[-1, 128, 64, 64]	131,200
InstanceNorm2d-4	[-1, 128, 64, 64]	0
LeakyReLU-5	[-1, 128, 64, 64]	0
Conv2d-6	[-1, 256, 32, 32]	524,544
InstanceNorm2d-7	[-1, 256, 32, 32]	0
LeakyReLU-8	[-1, 256, 32, 32]	0
Conv2d-9	[-1, 512, 16, 16]	2,097,664
InstanceNorm2d-10	[-1, 512, 16, 16]	0
LeakyReLU-11	[-1, 512, 16, 16]	0
Conv2d-12	[-1, 512, 8, 8]	4,194,816
InstanceNorm2d-13	[-1, 512, 8, 8]	0
LeakyReLU-14	[-1, 512, 8, 8]	0
Conv2d-15	[-1, 512, 4, 4]	4,194,816
InstanceNorm2d-16	[-1, 512, 4, 4]	0
LeakyReLU-17	[-1, 512, 4, 4]	0
Conv2d-18	[-1, 512, 2, 2]	4,194,816
InstanceNorm2d-19	[-1, 512, 2, 2]	0
LeakyReLU-20	[-1, 512, 2, 2]	0
Conv2d-21	[-1, 512, 1, 1]	4,194,816
LeakyReLU-22	[-1, 512, 1, 1]	0
ConvTranspose2d-23	[-1, 512, 2, 2]	4,194,816
InstanceNorm2d-24	[-1, 512, 2, 2]	0
Dropout-25	[-1, 512, 2, 2]	0
ReLU-26	[-1, 512, 2, 2]	0
ConvTranspose2d-27	[-1, 512, 4, 4]	8,389,120
InstanceNorm2d-28	[-1, 512, 4, 4]	0
Dropout-29	[-1, 512, 4, 4]	0
ReLU-30	[-1, 512, 4, 4]	0
ConvTranspose2d-31	[-1, 512, 8, 8]	8,389,120
InstanceNorm2d-32	[-1, 512, 8, 8]	0
Dropout-33	[-1, 512, 8, 8]	0
ReLU-34	[-1, 512, 8, 8]	0
ConvTranspose2d-35	[-1, 512, 16, 16]	8,389,120
InstanceNorm2d-36	[-1, 512, 16, 16]	0
Dropout-37	[-1, 512, 16, 16]	0
ReLU-38	[-1, 512, 16, 16]	0
ConvTranspose2d-39	[-1, 256, 32, 32]	4,194,560
InstanceNorm2d-40	[-1, 256, 32, 32]	0
Dropout-41	[-1, 256, 32, 32]	0
ReLU-42	[-1, 256, 32, 32]	0
ConvTranspose2d-43	[-1, 128, 64, 64]	1,048,704
InstanceNorm2d-44	[-1, 128, 64, 64]	0
Dropout-45	[-1, 128, 64, 64]	0
ReLU-46	[-1, 128, 64, 64]	0
ConvTranspose2d-47	[-1, 64, 128, 128]	262,208
InstanceNorm2d-48	[-1, 64, 128, 128]	0
Dropout-49	[-1, 64, 128, 128]	0
ReLU-50	[-1, 64, 128, 128]	0
ConvTranspose2d-51	[-1, 3, 256, 256]	6,147
Tanh-52	[-1, 3, 256, 256]	0

=====  
Total params: 54,409,603  
Trainable params: 54,409,603

## ▼ Data preparation

After the model is initialized, we then create a dataloader to sample the training data. In this paper, to sample the training data, you have to sequentially perform the following steps :

1. Randomly sample the data from the training set
2. Resizing both input and target to  $286 \times 286$ .
3. Randomly cropping back both images to size  $256 \times 256$ .
4. Random mirroring the images

TODO15: Implement a dataloader based on the description above. You are allowed to use the function in `torchvision.transforms` (<https://pytorch.org/vision/main/transforms.html>).

```
1 trainX.shape, trainY.shape
((400, 3, 256, 256), (400, 3, 256, 256))
```

```

1 # TODO15 implement a dataloader
2 from torch.utils.data import Dataset, DataLoader
3 from torchvision.transforms import v2
4
5 class P2P_dataset(Dataset):
6     def __init__(self, trainX, trainY):
7         trainX = torch.tensor(trainX).cuda()
8         trainY = torch.tensor(trainY).cuda()
9
10    self.train = torch.cat([trainX, trainY], 1)
11
12    self.transforms = v2.Compose([
13        v2.Resize((286, 286)),
14        v2.RandomResizedCrop(size=(256, 256)),
15        v2.RandomVerticalFlip(p=0.5),
16        v2.RandomHorizontalFlip(p=0.5),
17    ])
18
19    self.transformed_train = self.transforms(self.train)
20
21    def __len__(self):
22        return len(self.transformed_train)
23
24    def __getitem__(self, idx):
25        return self.transformed_train[idx, :3, :, :], self.transformed_train[idx, 3:, :, :]
26
27 p2p_dataset = P2P_dataset(trainX, trainY)
28 p2p_dataloader = DataLoader(p2p_dataset, batch_size=1, shuffle=True)

```

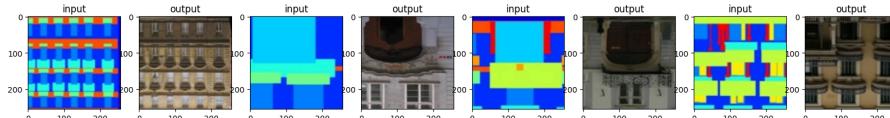
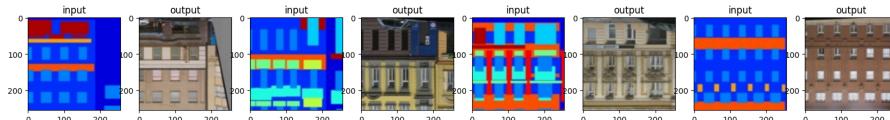
## ✓ Dataloader verification

TODO16: Show that the implemented dataloader is working as intended. For instance, are both input and output are flipped and cropped correctly? To obtain a full score, you have to show at least eight data points.

```

1 #TODO16 show that the dataloader is working properly
2 plt.figure(figsize=(20, 10))
3 for i in range(8):
4     x, y = next(iter(p2p_dataloader))
5     x = x.cpu().numpy().reshape(3, 256, 256)
6     y = y.cpu().numpy().reshape(3, 256, 256)
7     x = (0.5 * x.transpose((1, 2, 0)) + 0.5)[..., ::-1]
8     y = (0.5 * y.transpose((1, 2, 0)) + 0.5)[..., ::-1]
9     plt.subplot(2, 8, 2*i+1)
10    plt.title('input')
11    plt.imshow(x)
12    plt.subplot(2, 8, 2*i+2)
13    plt.title('output')
14    plt.imshow(y)

```



## ▼ Parameter Initialization

Model hyperparameters and optimizers have already been prepared.

```

1 import torch.optim as optim
2 from tqdm import tqdm
3 lr = 2e-4
4 LAMBDA = 100
5 BATCH_SIZE = 1
6 G_optimizer = optim.Adam(generator.parameters(), lr=lr, betas = (0.5, 0.999))
7 D_optimizer = optim.Adam(discriminator.parameters(), lr=lr, betas = (0.5, 0.999))
8
9 p2p_dataloader = DataLoader(p2p_dataset, batch_size=BATCH_SIZE, shuffle=True)

```

## ▼ Training loop

The training process has the following specific requirements:

- The objective is divided by 2 while optimizing  $D$ , which slows down the rate at which  $D$  learns relative to  $G$ .
- This paper trains the generator  $G$  to maximize  $\log D(x, G(x, z))$  instead of minimizing  $\log(1 - D(x, G(x, z)))$  as the latter term does not provide sufficient gradient.

TODO17: Sample the data using the dataloader.

TODO18: Calculate the discriminator loss and update the discriminator.

- During the update, the loss term  $\log(1 - D(x, G(x, z)))$  contains both generator and discriminator. However, we only want to update the discriminator. Therefore, you have to detach the input from the generator graph first. Read <https://pytorch.org/docs/stable/generated/torch.Tensor.detach.html> for additional detail.

TODO19: Calculate the generator loss and update the generator.

### HINT

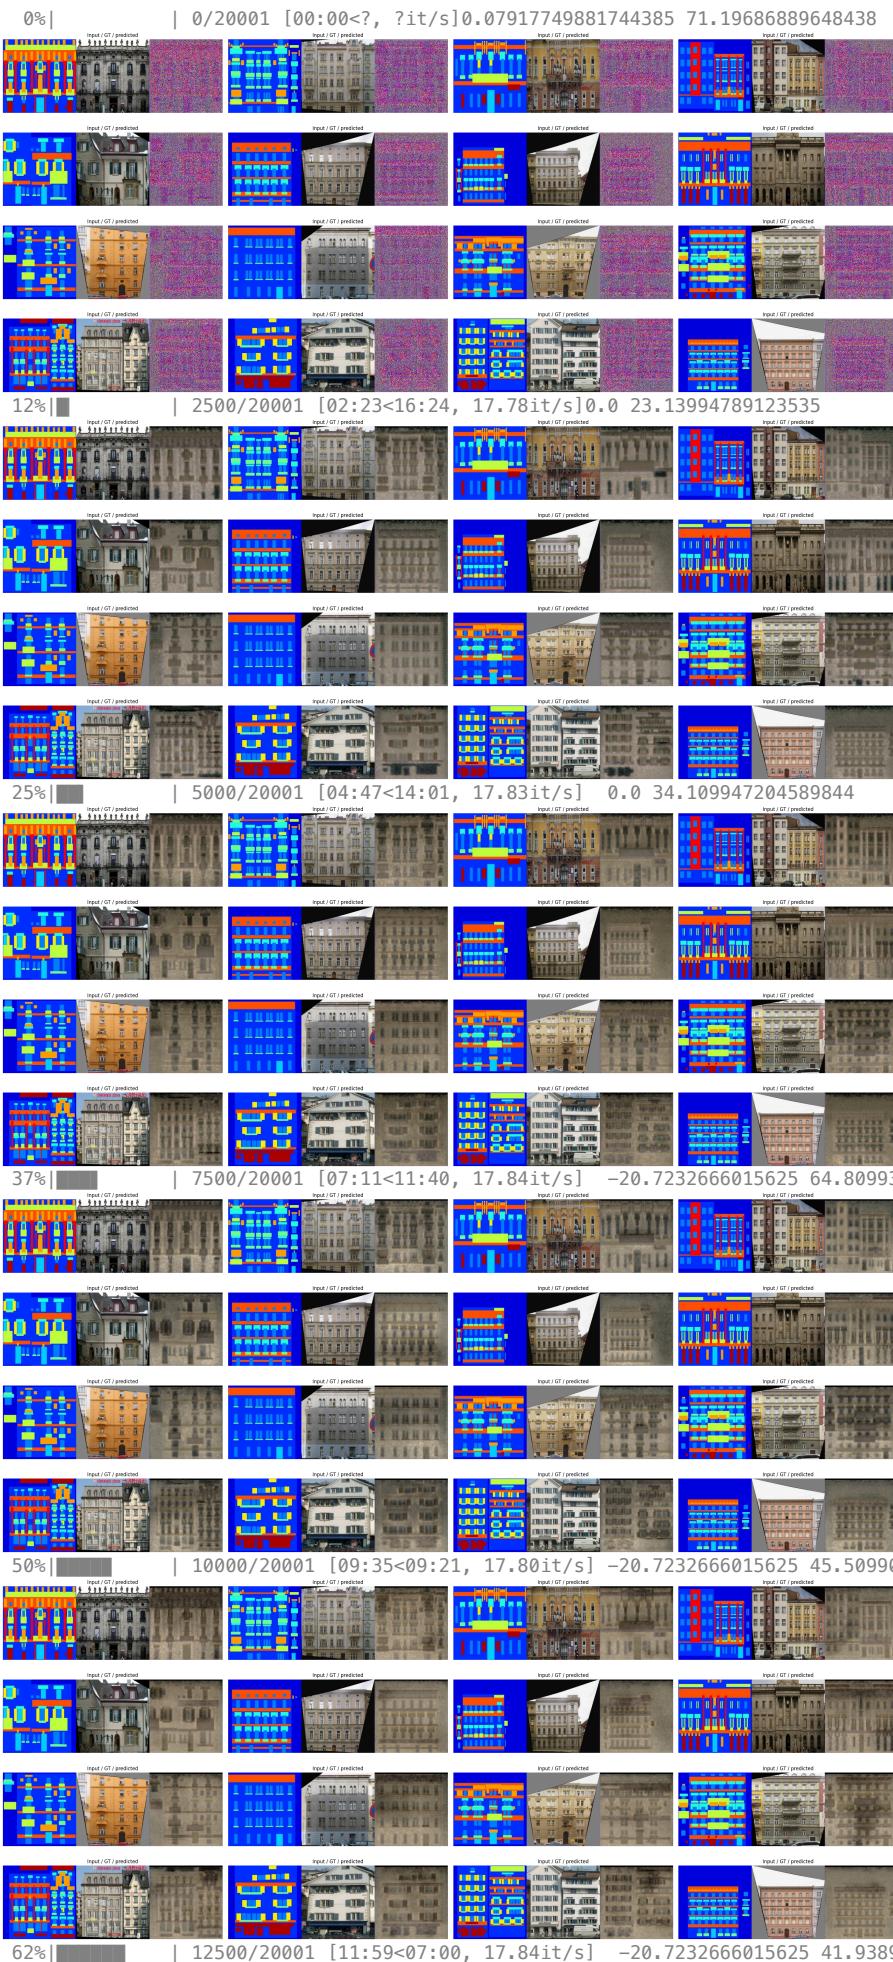
Hint 1: If you are struggling with this part, you could also read the PyTorch DCGAN tutorial as a guideline ([https://pytorch.org/tutorials/beginner/dcgan\\_faces\\_tutorial.html](https://pytorch.org/tutorials/beginner/dcgan_faces_tutorial.html)).

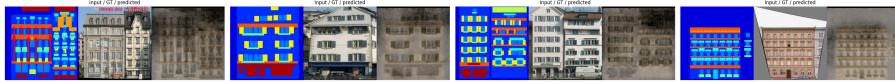
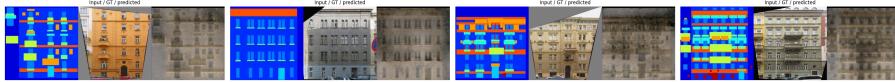
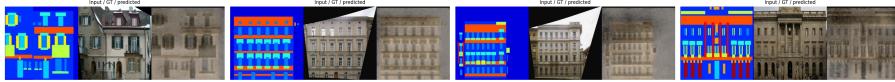
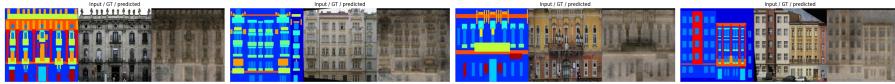
Hint 2: You could remove the L1 loss while debugging since the generator  $G$  could still generate the synthetic image even if the L1 loss is removed, though at the cost of increasing image artifacts.

#### Note

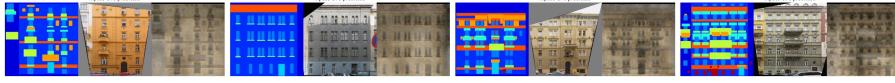
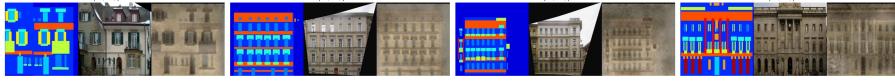
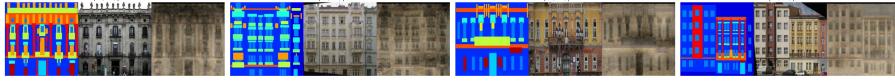
The training schedule in this homework is only one eighth of the original schedule. It is expected that the generated image quality is worse than the one shown in the paper. Nevertheless, the generated facade should still resemble an actual one.

```
1 losses = {'D' : [None], 'G' : [None]}
2
3 for i in tqdm(range(20001)):
4     #TOD017 sample the data from the dataloader
5     x, y = next(iter(p2p_dataloader))
6     #TOD018 calculate the discriminator loss and update the discriminator
7     g_x = generator(x)
8     if i%1 == 0:
9         f_pred = discriminator(torch.cat([x, g_x.detach()], 1)) + 1e-9
10        t_pred = discriminator(torch.cat([x, y], 1)) + 1e-9
11    #     print(f"true pred : {t_pred.item()}, false pred : {f_pred.item()}")
12        d_loss = torch.log(f_pred)-torch.log(t_pred)
13        losses['D'].append(d_loss.item())
14
15        D_optimizer.zero_grad()
16        d_loss.backward()
17        D_optimizer.step()
18
19 #TOD019 calculate the generator loss and update the generator
20 g_loss = torch.log(discriminator(torch.cat([x, y], 1)) + 1e-9) - torch.log(discriminator(torch.cat([x, g_x], 1)) + 1e-9)
21 losses['G'].append(g_loss.item())
22
23 G_optimizer.zero_grad()
24 g_loss.backward()
25 G_optimizer.step()
26
27 # Output visualization : If your reimplementation is correct, the generated images should start resembling a facade after
28 if(i % 2500 == 0):
29     with torch.no_grad():
30         print(losses['D'][-1], losses['G'][-1])
31         plt.figure(figsize = (40, 16))
32         gs1 = gridspec.GridSpec(4, 4)
33         gs1.update(wspace=0.025)
34
35         sampleX_vis = 0.5 * valX[:16][:, ::-1, :, :] + 0.5
36         sampleY = 0.5 * valY[:16][:, ::-1, :, :] + 0.5
37         sampleX = torch.tensor(valX[:16]).cuda()
38         pred_val = 0.5 * generator(sampleX).cpu().detach().numpy()[:, ::-1, :, :] + 0.5
39         vis = np.concatenate([sampleX_vis, sampleY, pred_val], axis = 3)
40         for i in range(vis.shape[0]):
41             ax1 = plt.subplot(gs1[i])
42             plt.title('Input / GT / predicted')
43             plt.axis('off')
44             plt.imshow( vis[i].transpose(1, 2, 0) )
45             plt.show()
```

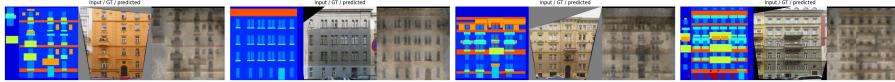
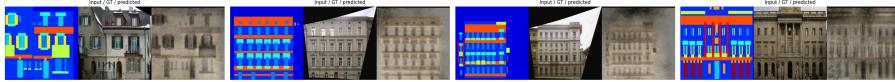
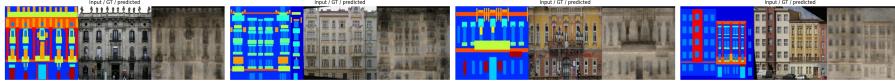




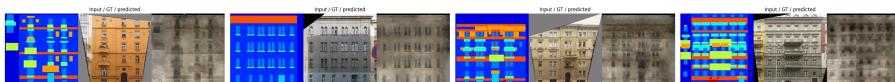
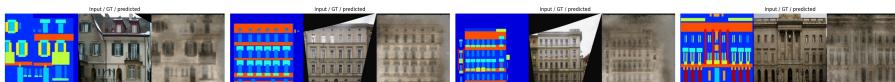
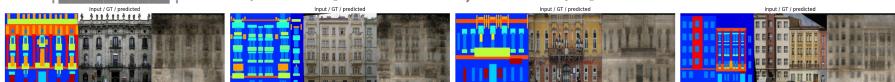
75% | 15000/20001 [14:23<04:41, 17.79it/s] -20.7232666015625 38.8269



87% | 17500/20001 [16:48<02:20, 17.82it/s]-20.7232666015625 38.63603



100% | 20000/20001 [19:12<00:00, 17.80it/s]-20.7232666015625 43.405570



100% | 20001/20001 [19:16<00:00, 17.29it/s]

